

lab02

November 7, 2025

Nº2

8 -303 -23

```
[60]: %pip install catboost
```

```
Requirement already satisfied: catboost in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (1.2.8)
Requirement already satisfied: graphviz in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from catboost) (0.21)
Requirement already satisfied: matplotlib in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from catboost) (3.10.7)
Requirement already satisfied: numpy<3.0,>=1.16.0 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from catboost) (2.3.4)
Requirement already satisfied: pandas>=0.24 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from catboost) (2.3.3)
Requirement already satisfied: scipy in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from catboost) (1.16.3)
Requirement already satisfied: plotly in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from catboost) (6.4.0)
Requirement already satisfied: six in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from catboost) (1.17.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from pandas>=0.24->catboost) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from pandas>=0.24->catboost) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from pandas>=0.24->catboost) (2025.2)
Requirement already satisfied: contourpy>=1.0.1 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from matplotlib->catboost) (1.3.3)
Requirement already satisfied: cycler>=0.10 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from matplotlib->catboost) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from matplotlib->catboost) (4.60.1)
Requirement already satisfied: kiwisolver>=1.3.1 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from matplotlib->catboost) (1.4.9)
Requirement already satisfied: packaging>=20.0 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from matplotlib->catboost) (25.0)
```

```
Requirement already satisfied: pillow>=8 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from matplotlib->catboost) (12.0.0)
Requirement already satisfied: pyparsing>=3 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from matplotlib->catboost) (3.2.5)
Requirement already satisfied: narwhals>=1.15.1 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from plotly->catboost) (2.10.2)
Note: you may need to restart the kernel to use updated packages.
```

```
[61]: import os
from pathlib import Path
import math
import warnings
warnings.filterwarnings("ignore")

import numpy as np
import pandas as pd
import seaborn as sns

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV, StratifiedKFold
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix

from sklearn.metrics import (
    accuracy_score, precision_score, recall_score,
    f1_score, roc_auc_score
)

from catboost import CatBoostClassifier, Pool, cv
```

```
[62]: import os
from pathlib import Path

IN_COLAB = "COLAB_GPU" in os.environ or "google.colab" in str(get_ipython())

if not IN_COLAB:
    ROOT = Path.cwd()

    if ROOT.name == "lab02":
```

```

    LAB_ROOT = ROOT
elif (ROOT / "lab02").exists():
    LAB_ROOT = ROOT / "lab02"
elif ROOT.name == "notebooks":
    LAB_ROOT = ROOT.parent
else:
    raise FileNotFoundError("          lab02          .")

DATA_RAW = LAB_ROOT / "data" / "raw"
DATA_RAW.mkdir(parents=True, exist_ok=True)

csv_path = DATA_RAW / "telco.csv"

else:
    from urllib.request import urlretrieve

    url = "https://raw.githubusercontent.com/OlegTertychnyi/ML-labs-Tert/main/
↪lab02/data/raw/telco.csv"
    csv_path = Path("/content/telco.csv")

    if not csv_path.exists():
        urlretrieve(url, csv_path)
        print(" Dataset downloaded to Colab")

csv_path

```

[62]: PosixPath('/home/optert/ML-labs-Tert/lab02/data/raw/telco.csv')

[63]: df = pd.read_csv(csv_path)
df

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	\
0	7590-VHVEG	Female	0	Yes	No	1	
1	5575-GNVDE	Male	0	No	No	34	
2	3668-QPYBK	Male	0	No	No	2	
3	7795-CFOCW	Male	0	No	No	45	
4	9237-HQITU	Female	0	No	No	2	
...	
7038	6840-RESVB	Male	0	Yes	Yes	24	
7039	2234-XADUH	Female	0	Yes	Yes	72	
7040	4801-JZAZL	Female	0	Yes	Yes	11	
7041	8361-LTMKD	Male	1	Yes	No	4	
7042	3186-AJIEK	Male	0	No	No	66	

PhoneService MultipleLines InternetService OnlineSecurity ... \

0	No	No phone service	DSL	No	...
1	Yes	No	DSL	Yes	...
2	Yes	No	DSL	Yes	...
3	No	No phone service	DSL	Yes	...
4	Yes	No	Fiber optic	No	...
...
7038	Yes	Yes	DSL	Yes	...
7039	Yes	Yes	Fiber optic	No	...
7040	No	No phone service	DSL	Yes	...
7041	Yes	Yes	Fiber optic	No	...
7042	Yes	No	Fiber optic	Yes	...
DeviceProtection TechSupport StreamingTV StreamingMovies Contract \					
0	No	No	No	No	Month-to-month
1	Yes	No	No	No	One year
2	No	No	No	No	Month-to-month
3	Yes	Yes	No	No	One year
4	No	No	No	No	Month-to-month
...
7038	Yes	Yes	Yes	Yes	One year
7039	Yes	No	Yes	Yes	One year
7040	No	No	No	No	Month-to-month
7041	No	No	No	No	Month-to-month
7042	Yes	Yes	Yes	Yes	Two year
PaperlessBilling PaymentMethod MonthlyCharges TotalCharges \					
0	Yes	Electronic check	29.85	29.85	
1	No	Mailed check	56.95	1889.5	
2	Yes	Mailed check	53.85	108.15	
3	No	Bank transfer (automatic)	42.30	1840.75	
4	Yes	Electronic check	70.70	151.65	
...
7038	Yes	Mailed check	84.80	1990.5	
7039	Yes	Credit card (automatic)	103.20	7362.9	
7040	Yes	Electronic check	29.60	346.45	
7041	Yes	Mailed check	74.40	306.6	
7042	Yes	Bank transfer (automatic)	105.65	6844.5	
Churn					
0	No				
1	No				
2	Yes				
3	No				
4	Yes				
...	...				
7038	No				
7039	No				

7040	No
7041	Yes
7042	No

[7043 rows x 21 columns]

```
[64]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   customerID      7043 non-null    object  
 1   gender          7043 non-null    object  
 2   SeniorCitizen   7043 non-null    int64  
 3   Partner         7043 non-null    object  
 4   Dependents     7043 non-null    object  
 5   tenure          7043 non-null    int64  
 6   PhoneService    7043 non-null    object  
 7   MultipleLines   7043 non-null    object  
 8   InternetService 7043 non-null    object  
 9   OnlineSecurity  7043 non-null    object  
 10  OnlineBackup    7043 non-null    object  
 11  DeviceProtection 7043 non-null    object  
 12  TechSupport    7043 non-null    object  
 13  StreamingTV    7043 non-null    object  
 14  StreamingMovies 7043 non-null    object  
 15  Contract        7043 non-null    object  
 16  PaperlessBilling 7043 non-null    object  
 17  PaymentMethod   7043 non-null    object  
 18  MonthlyCharges 7043 non-null    float64 
 19  TotalCharges    7043 non-null    object  
 20  Churn           7043 non-null    object  
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
[65]: for col in df.select_dtypes(include=['object']).columns:  
    df[col] = df[col].astype(str).str.strip()  
df["TotalCharges"] = pd.to_numeric(df["TotalCharges"], errors="coerce")  
df["Churn"] = (df["Churn"] == "Yes").astype(int)  
df["TotalCharges"] = df["TotalCharges"].fillna(df["TotalCharges"].median())
```

```
[66]: print("      : ", df.shape)
df.info()
```

```
: (7043, 21)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   customerID        7043 non-null   object  
 1   gender             7043 non-null   object  
 2   SeniorCitizen     7043 non-null   int64  
 3   Partner            7043 non-null   object  
 4   Dependents         7043 non-null   object  
 5   tenure              7043 non-null   int64  
 6   PhoneService       7043 non-null   object  
 7   MultipleLines      7043 non-null   object  
 8   InternetService    7043 non-null   object  
 9   OnlineSecurity     7043 non-null   object  
 10  OnlineBackup        7043 non-null   object  
 11  DeviceProtection   7043 non-null   object  
 12  TechSupport         7043 non-null   object  
 13  StreamingTV         7043 non-null   object  
 14  StreamingMovies     7043 non-null   object  
 15  Contract            7043 non-null   object  
 16  PaperlessBilling   7043 non-null   object  
 17  PaymentMethod       7043 non-null   object  
 18  MonthlyCharges     7043 non-null   float64 
 19  TotalCharges        7043 non-null   float64 
 20  Churn               7043 non-null   int64  
dtypes: float64(2), int64(3), object(16)
memory usage: 1.1+ MB
```

target

```
[67]: TARGET = "Churn"
ID_COLS = ["customerID"] if "customerID" in df.columns else []

X = df.drop(columns=[TARGET] + ID_COLS)

num_cols = X.select_dtypes(include=[np.number]).columns.tolist()
cat_cols = X.select_dtypes(include=["object"]).columns.tolist()

num_cols, cat_cols
```

```
[67]: ([['SeniorCitizen', 'tenure', 'MonthlyCharges', 'TotalCharges'],
      ['gender'],
      'Partner',
```

```
'Dependents',
'PhoneService',
'MultipleLines',
'InternetService',
'OnlineSecurity',
'OnlineBackup',
'DeviceProtection',
'TechSupport',
'StreamingTV',
'StreamingMovies',
'Contract',
'PaperlessBilling',
'PaymentMethod'])
```

```
[68]: df.isna().sum().sort_values(ascending=False)
```

```
[68]: customerID      0
gender          0
SeniorCitizen    0
Partner          0
Dependents       0
tenure           0
PhoneService     0
MultipleLines     0
InternetService   0
OnlineSecurity    0
OnlineBackup       0
DeviceProtection  0
TechSupport        0
StreamingTV       0
StreamingMovies    0
Contract          0
PaperlessBilling   0
PaymentMethod      0
MonthlyCharges     0
TotalCharges       0
Churn             0
dtype: int64
```

```
[69]: df[num_cols].describe().T
```

```
[69]:      count      mean       std      min      25%      50%  \
SeniorCitizen  7043.0  0.162147  0.368612  0.00  0.000  0.000
tenure         7043.0  32.371149  24.559481  0.00  9.000  29.000
```

```

MonthlyCharges    7043.0      64.761692      30.090047   18.25     35.500     70.350
TotalCharges      7043.0    2281.916928    2265.270398   18.80     402.225   1397.475

                           75%      max
SeniorCitizen        0.00     1.00
tenure                55.00    72.00
MonthlyCharges       89.85   118.75
TotalCharges        3786.60  8684.80

```

[70]: df[cat_cols].describe().T

	count	unique	top	freq
gender	7043	2	Male	3555
Partner	7043	2	No	3641
Dependents	7043	2	No	4933
PhoneService	7043	2	Yes	6361
MultipleLines	7043	3	No	3390
InternetService	7043	3	Fiber optic	3096
OnlineSecurity	7043	3	No	3498
OnlineBackup	7043	3	No	3088
DeviceProtection	7043	3	No	3095
TechSupport	7043	3	No	3473
StreamingTV	7043	3	No	2810
StreamingMovies	7043	3	No	2785
Contract	7043	3	Month-to-month	3875
PaperlessBilling	7043	2	Yes	4171
PaymentMethod	7043	4	Electronic check	2365

. Monthly-
Charges . «No»
«Month-to-month», .

```

[71]: n = len(num_cols)
rows = (n + 1) // 2

fig, axes = plt.subplots(rows, 2, figsize=(12, 4 * rows))
axes = axes.flatten()

for i, col in enumerate(num_cols):
    axes[i].hist(df[col], bins=30)
    axes[i].set_title(f" {col}")
    axes[i].set_xlabel(col)
    axes[i].set_ylabel(" ")
    axes[i].grid(alpha=0.3)

for j in range(i+1, len(axes)):


```

```

fig.delaxes(axes[j])

plt.tight_layout()
plt.show()

```



SeniorCitizen : tenure,
 MonthlyCharges ,
 TotalCharges , . . .

```

[72]: n_cols = 3
n_rows = math.ceil(len(cat_cols) / n_cols)

fig, axes = plt.subplots(n_rows, n_cols, figsize=(5 * n_cols, 4 * n_rows))
axes = axes.flatten()

for ax, col in zip(axes, cat_cols):
    sns.countplot(data=df, x=col, ax=ax)
    ax.set_title(col)
    ax.tick_params(axis="x", rotation=45, labelrotation=45)
    ax.grid(axis="y", alpha=0.3)

for ax in axes[len(cat_cols):]:


```

```

fig.delaxes(ax)

plt.tight_layout()
plt.show()

```




,
 TechSupport DeviceProtection.
 .
 ,

«No» —
 — «Month-to-month»,
 — Electronic check.

OnlineSecurity,

```
[74]: sns.pairplot(df[num_cols + ['Churn']], hue='Churn', plot_kws={'alpha': 0.5})
plt.show()
```



,
 .
 ,
 churn,

```
[75]: n_cols = 3
n_rows = math.ceil(len(num_cols) / n_cols)
```

```


fig, axes = plt.subplots(n_rows, n_cols, figsize=(5 * n_cols, 4 * n_rows))
axes = axes.flatten()

for ax, col in zip(axes, num_cols):
    sns.boxplot(data=df, x="Churn", y=col, ax=ax)
    ax.set_title(f"{col} vs Churn")
    ax.set_xlabel("Churn")
    ax.set_ylabel(col)
    ax.grid(axis="y", alpha=0.3)

for ax in axes[len(num_cols):]:
    fig.delaxes(ax)

plt.tight_layout()
plt.show()

```



churn

tenure

MonthlyCharges,
TotalCharges,

,

[76]: corr = df[num_cols + ["Churn"]].corr()

```


plt.figure(figsize=(8, 6))
sns.heatmap(
    corr,
    annot=True,
    cmap="coolwarm",
    square=True,

```

```

    cbar=True
)
plt.title(" ")
plt.show()

```



, tenure TotalCharges ,
Churn tenure — .
, .
train/test- (20%).
proc_trees), SVM/KNN — : customerID, Churn,
Boost : — One-Hot (pre-
ing_rate, iterations) : StandardScaler (preproc_svm_knn). Cat-
(cat_features_idx), (depth, learn-
5-fold - - ; Pool.

```
[77]: TARGET = "Churn"
ID_COLS = ["customerID"] if "customerID" in df.columns else []
```

```
X = df.drop(columns=[TARGET] + ID_COLS)
y = df[TARGET]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

```
[78]: num_cols = X.select_dtypes(include=["number"]).columns.tolist()
cat_cols = X.select_dtypes(include=["object"]).columns.tolist()

num_pipe = Pipeline([
    ("imputer", SimpleImputer(strategy="median"))
])

cat_pipe = Pipeline([
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

preproc_trees = ColumnTransformer(
    transformers=[
        ("num", num_pipe, num_cols),
        ("cat", cat_pipe, cat_cols)
    ]
)
```

```
[79]: num_pipe_scaled = Pipeline([
    ("imputer", SimpleImputer(strategy="median")),
    ("scaler", StandardScaler())
])

cat_pipe = Pipeline([
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

preproc_svm_knn = ColumnTransformer(
    transformers=[
        ("num", num_pipe_scaled, num_cols),
        ("cat", cat_pipe, cat_cols)
    ]
)
```

```
[80]: cat_features_idx = [X.columns.get_loc(c) for c in cat_cols]
```

```
[81]: param_grid = [
    {"depth": 4, "learning_rate": 0.1, "iterations": 600},
    {"depth": 6, "learning_rate": 0.05, "iterations": 800},
    {"depth": 8, "learning_rate": 0.03, "iterations": 1000},
]
```

```
[82]: skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
```

```
[83]: train_pool = Pool(X_train, y_train, cat_features=cat_features_idx)
```

CV	CatBoost	-	.	5-fold
	ROC-AUC			
		.		

```
[84]: results = []

for params in param_grid:
    print(f"          : {params}")

    cv_results = cv(
        params={
            "loss_function": "Logloss",
            "eval_metric": "AUC",
            "random_seed": 42,
            "early_stopping_rounds": 50,
            **params
        },
        pool=train_pool,
        fold_count=5,
        shuffle=True,
        partition_random_seed=42,
        verbose=False
    )

    best_auc = cv_results["test-AUC-mean"].max()
    best_iter = cv_results["test-AUC-mean"].idxmax()

    results.append({
        "params": params,
        "best_auc": best_auc,
        "best_iter": best_iter
    })

results_df = pd.DataFrame(results)
results_df
```

: {'depth': 4, 'learning_rate': 0.1, 'iterations': 600}

Training on fold [0/5]

```
bestTest = 0.8684847236
bestIteration = 106

Training on fold [1/5]

bestTest = 0.8321700354
bestIteration = 148

Training on fold [2/5]

bestTest = 0.851707786
bestIteration = 88

Training on fold [3/5]

bestTest = 0.8548361689
bestIteration = 70

Training on fold [4/5]

bestTest = 0.8440347308
bestIteration = 260

: {'depth': 6, 'learning_rate': 0.05, 'iterations': 800}
Training on fold [0/5]

bestTest = 0.8670164639
bestIteration = 132

Training on fold [1/5]

bestTest = 0.8289507699
bestIteration = 77

Training on fold [2/5]

bestTest = 0.8498093484
bestIteration = 97

Training on fold [3/5]

bestTest = 0.8553935825
bestIteration = 51

Training on fold [4/5]

bestTest = 0.842360468
```

```
bestIteration = 259

        : {'depth': 8, 'learning_rate': 0.03, 'iterations': 1000}
Training on fold [0/5]

bestTest = 0.8671497585
bestIteration = 179

Training on fold [1/5]

bestTest = 0.8309259529
bestIteration = 143

Training on fold [2/5]

bestTest = 0.8521965327
bestIteration = 117

Training on fold [3/5]

bestTest = 0.8549432084
bestIteration = 93

Training on fold [4/5]

bestTest = 0.8393698463
bestIteration = 280
```

		params	best_auc	best_iter
0	{'depth': 4, 'learning_rate': 0.1, 'iterations...}	0.849032	148	
1	{'depth': 6, 'learning_rate': 0.05, 'iteration...}	0.848564	259	
2	{'depth': 8, 'learning_rate': 0.03, 'iteration...}	0.848585	142	

CatBoost, iterations, Pool.

```
[85]: best_cfg = results_df.iloc[results_df["best_auc"].idxmax()]
      best_cfg
```

```
[85]: params      {'depth': 4, 'learning_rate': 0.1, 'iterations...  
       best_auc          0.849032  
       best_iter         148  
       Name: 0, dtype: object
```

```
[86]: best_params = best_cfg["params"].copy()
       best_params.pop("iterations")
       best_iter = int(best_cfg["best_iter"])
```

```

model_cat = CatBoostClassifier(
    loss_function="Logloss",
    eval_metric="AUC",
    random_seed=42,
    iterations=best_iter,
    **best_params
)

model_cat.fit(train_pool, verbose=False)

```

[86]: <catboost.core.CatBoostClassifier at 0x7fc598572e10>

- GridSearchCV

```

[87]: rf_param_grid = {
    "model__n_estimators": [200, 400, 800],
    "model__max_depth": [None, 10, 15],
    "model__min_samples_leaf": [1, 3]
}

```

```

[88]: rf_pipeline = Pipeline([
    ("preprocess", preproc_trees),
    ("model", RandomForestClassifier(
        class_weight="balanced",
        random_state=42,
        n_jobs=-1
    ))
])

```

[89]: skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

```

[90]: rf_grid = GridSearchCV(
    estimator=rf_pipeline,
    param_grid=rf_param_grid,
    scoring="roc_auc",
    cv=skf,
    n_jobs=-1,
    verbose=1
)

rf_grid.fit(X_train, y_train)

```

Fitting 5 folds for each of 18 candidates, totalling 90 fits

```

[90]: GridSearchCV(cv=StratifiedKFold(n_splits=5, random_state=42, shuffle=True),
                  estimator=Pipeline(steps=[('preprocess',

```

```

Pipeline(steps=[('imputer',
    SimpleImputer(strategy='median'))],
['SeniorCitizen',
'tenure',
'MonthlyCharges',
'TotalCharges']),
('cat',
Pipeline(steps=[('imputer',
    SimpleImputer(strategy='most_frequent')),...
'DeviceProtection',
'TechSupport',
'StreamingTV',
'StreamingMovies',
'Contract',
'PaperlessBilling',
'PaymentMethod')]]),
('model',
RandomForestClassifier(class_weight='balanced',
n_jobs=-1,
random_state=42))],
n_jobs=-1,
param_grid={'model__max_depth': [None, 10, 15],
'model__min_samples_leaf': [1, 3],
'model__n_estimators': [200, 400, 800]},
scoring='roc_auc', verbose=1)

```

```
[91]: best_rf = rf_grid.best_estimator_
best_rf_params = rf_grid.best_params_
best_rf_score = rf_grid.best_score_

best_rf_params, best_rf_score
```

```
[91]: ({'model__max_depth': 10,
'model__min_samples_leaf': 3,
'model__n_estimators': 800},
np.float64(0.8452849518294346))
```

SVM kNN

```
[92]: svc_param_grid = {
    "model__C": [0.5, 1, 2, 4],
    "model__kernel": ["rbf"],
}
```

```
[93]: svc_pipeline = Pipeline([
    ("preprocess", preproc_svm_knn),
    ("model", SVC(
```

```
        probability=True,
        class_weight="balanced",
        random_state=42
    ))
])
])
```

```
[94]: skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
```

```
[95]: svc_grid = GridSearchCV(
    estimator=svc_pipeline,
    param_grid=svc_param_grid,
    scoring="roc_auc",
    cv=skf,
    n_jobs=-1,
    verbose=1
)

svc_grid.fit(X_train, y_train)
```

Fitting 5 folds for each of 4 candidates, totalling 20 fits

```
[95]: GridSearchCV(cv=StratifiedKFold(n_splits=5, random_state=42, shuffle=True),
                    estimator=Pipeline(steps=[('preprocess',
                                              ColumnTransformer(transformers=[('num',
                                                                 Pipeline(steps=[('imputer',
                                                                 SimpleImputer(strategy='median'))),
                                                                 ('scaler',
                                                                 StandardScaler()))]),
                                              ['SeniorCitizen',
                                               'tenure',
                                               'MonthlyCharges',
                                               'TotalCharges']),
                                              ('cat',
                                               Pipeline(steps=[('imputer',
                                               SimpleImputer...
                                               'MultipleLines',
                                               'InternetService',
                                               'OnlineSecurity',
                                               'OnlineBackup',
                                               'DeviceProtection',
                                               'TechSupport',
                                               'StreamingTV',
                                               'StreamingMovies',
                                               'Contract',
                                               'PaperlessBilling',
                                               'PaymentMethod')])),
                                              ('model',
```

```

        SVC(class_weight='balanced',
             probability=True,
             random_state=42))),

n_jobs=-1,
param_grid={'model__C': [0.5, 1, 2, 4], 'model__kernel': ['rbf']},
scoring='roc_auc', verbose=1)

```

```
[96]: best_svc = svc_grid.best_estimator_
best_svc_params = svc_grid.best_params_
best_svc_score = svc_grid.best_score_

best_svc_params, best_svc_score
```

```
[96]: ({'model__C': 0.5, 'model__kernel': 'rbf'}, np.float64(0.8322140234176816))
```

```
[97]: knn_param_grid = {
    "model__n_neighbors": [5, 9, 15, 25],
    "model__weights": ["uniform", "distance"],
    "model__p": [1, 2]
}
```

```
[98]: knn_pipeline = Pipeline([
    ("preprocess", preproc_svm_knn),
    ("model", KNeighborsClassifier())
])
```

```
[99]: skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
```

```
[100]: knn_grid = GridSearchCV(
    estimator=knn_pipeline,
    param_grid=knn_param_grid,
    scoring="roc_auc",
    cv=skf,
    n_jobs=-1,
    verbose=1
)

knn_grid.fit(X_train, y_train)
```

Fitting 5 folds for each of 16 candidates, totalling 80 fits

```
[100]: GridSearchCV(cv=StratifiedKFold(n_splits=5, random_state=42, shuffle=True),
                    estimator=Pipeline(steps=[('preprocess',
                                              ColumnTransformer(transformers=[('num',

```

```
Pipeline(steps=[('imputer',
                  SimpleImputer(strategy='median')),
                  ('scaler',
```

```

        StandardScaler())]),
['SeniorCitizen',
'tenure',
'MonthlyCharges',
'TotalCharges']),
('cat',
Pipeline(steps=[('imputer',
SimpleImputer..
'MultipleLines',
'InternetService',
'OnlineSecurity',
'OnlineBackup',
'DeviceProtection',
'TechSupport',
'StreamingTV',
'StreamingMovies',
'Contract',
'PaperlessBilling',
'PaymentMethod')]]),
('model', KNeighborsClassifier())),
n_jobs=-1,
param_grid={'model__n_neighbors': [5, 9, 15, 25],
            'model__p': [1, 2],
            'model__weights': ['uniform', 'distance']},
scoring='roc_auc', verbose=1)

```

```
[101]: best_knn = knn_grid.best_estimator_
best_knn_params = knn_grid.best_params_
best_knn_score = knn_grid.best_score_

best_knn_params, best_knn_score
```

```
[101]: {'model__n_neighbors': 25, 'model__p': 2, 'model__weights': 'uniform'},
np.float64(0.8319873215684025)
```

```
[102]: def evaluate_model(name, model, X_test, y_test, proba=True):
    if name == "CatBoost":
        y_pred = model.predict(Pool(X_test, cat_features=cat_features_idx))
        y_proba = model.predict_proba(Pool(X_test, ↴
                                         cat_features=cat_features_idx))[:, 1]
    else:
        y_pred = model.predict(X_test)
        y_proba = model.predict_proba(X_test)[:, 1] if proba else None

    return {
```

```

    "model": name,
    "accuracy": accuracy_score(y_test, y_pred),
    "precision": precision_score(y_test, y_pred),
    "recall": recall_score(y_test, y_pred),
    "f1": f1_score(y_test, y_pred),
    "roc_auc": roc_auc_score(y_test, y_proba)
}

```

```
[103]: results = []

results.append(evaluate_model("CatBoost", model_cat, X_test, y_test))
results.append(evaluate_model("RandomForest", best_rf, X_test, y_test))
results.append(evaluate_model("SVC", best_svc, X_test, y_test))
results.append(evaluate_model("KNN", best_knn, X_test, y_test))

df_results = pd.DataFrame(results)
df_results
```

	model	accuracy	precision	recall	f1	roc_auc
0	CatBoost	0.804826	0.668942	0.524064	0.587706	0.846785
1	RandomForest	0.767921	0.546906	0.732620	0.626286	0.841058
2	SVC	0.743790	0.511304	0.786096	0.619600	0.825052
3	KNN	0.784954	0.598886	0.574866	0.586630	0.827011

. CatBoost , ROC-AUC precision,
 RandomForest SVC recall. KNN , .

```
[104]: import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, auc

plt.figure(figsize=(8, 6))

y_proba = model_cat.predict_proba(Pool(X_test, cat_features=cat_features_idx))[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_proba)
plt.plot(fpr, tpr, label=f"CatBoost (AUC = {auc(fpr, tpr):.3f})")

y_proba = best_rf.predict_proba(X_test)[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_proba)
plt.plot(fpr, tpr, label=f"RandomForest (AUC = {auc(fpr, tpr):.3f})")

y_proba = best_svc.predict_proba(X_test)[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_proba)
plt.plot(fpr, tpr, label=f"SVC (AUC = {auc(fpr, tpr):.3f})")


y_proba = best_knn.predict_proba(X_test)[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_proba)
```

```

plt.plot(fpr, tpr, label=f"KNN (AUC = {auc(fpr, tpr):.3f})")

plt.plot([0, 1], [0, 1], "k--", alpha=0.5)
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC-")
plt.legend()
plt.grid(alpha=0.3)
plt.show()

```



ROC- , CatBoost , RandomForest , . SVC KNN
, .
Precision-matrix

```
[105]: def plot_cm(ax, y_true, y_pred, title, normalize=True):
    cm = confusion_matrix(y_true, y_pred, labels=[0,1])
    if normalize:
        cm = cm.astype(float) / cm.sum(axis=1, keepdims=True)
    sns.heatmap(cm, annot=True, fmt=".2f" if normalize else "d",
```

```

        cmap="Blues", cbar=False, ax=ax,
        xticklabels=["Pred 0", "Pred 1"], yticklabels=["True 0", "True 1"])
    ax.set_title(title)
    ax.set_xlabel("")
    ax.set_ylabel("")


y_pred_cat = model_cat.predict(Pool(X_test, cat_features=cat_features_idx))
y_pred_rf = best_rf.predict(X_test)
y_pred_svc = best_svc.predict(X_test)
y_pred_knn = best_knn.predict(X_test)

fig, axes = plt.subplots(2, 2, figsize=(10, 8))
axes = axes.ravel()

plot_cm(axes[0], y_test, y_pred_cat, "CatBoost (CM, row-normalized)")
plot_cm(axes[1], y_test, y_pred_rf, "RandomForest (CM, row-normalized)")
plot_cm(axes[2], y_test, y_pred_svc, "SVC (CM, row-normalized)")
plot_cm(axes[3], y_test, y_pred_knn, "KNN (CM, row-normalized)")

plt.tight_layout()
plt.show()

```



RandomForest SVC ,
 (True 0),
 precision recall .
 heatmap


```
[106]: df_hm = df_results.copy()
df_hm = df_hm.set_index("model")

plt.figure(figsize=(8, 4.5))

sns.heatmap(
    df_hm,
    annot=True,
    fmt=' .5g',
    cmap="coolwarm",
    linewidths=.5,
    cbar=True
)

plt.tight_layout()
```

```
plt.show()
```



CatBoost
SVC
, (ROC-AUC, precision,
RandomForest
CatBoost.
, recall).

CatBoost
, Telco Customer Churn
One-Hot
, RandomForest, SVC, KNN, CatBoost
target-based (ordinal / target statistics),

, CatBoost
,
Ordered Boosting ,