

lab03

November 28, 2025

```
[ ]: %pip install catboost optuna optuna-integration[catboost]
```

```
Requirement already satisfied: catboost in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (1.2.8)
Requirement already satisfied: optuna in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (4.6.0)
Requirement already satisfied: optuna-integration[catboost] in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (4.6.0)
Requirement already satisfied: graphviz in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from catboost) (0.21)
Requirement already satisfied: matplotlib in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from catboost) (3.10.7)
Requirement already satisfied: numpy<3.0,>=1.16.0 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from catboost) (1.26.4)
Requirement already satisfied: pandas>=0.24 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from catboost) (2.3.3)
Requirement already satisfied: scipy in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from catboost) (1.16.3)
Requirement already satisfied: plotly in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from catboost) (6.4.0)
Requirement already satisfied: six in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from catboost) (1.17.0)
Requirement already satisfied: alembic>=1.5.0 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from optuna) (1.17.1)
Requirement already satisfied: colorlog in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from optuna) (6.10.1)
Requirement already satisfied: packaging>=20.0 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from optuna) (25.0)
Requirement already satisfied: sqlalchemy>=1.4.2 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from optuna) (2.0.44)
Requirement already satisfied: tqdm in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from optuna) (4.67.1)
Requirement already satisfied: PyYAML in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from optuna) (6.0.3)
Requirement already satisfied: Mako in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from alembic>=1.5.0->optuna) (1.3.10)
Requirement already satisfied: typing-extensions>=4.12 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from alembic>=1.5.0->optuna) (4.15.0)
```

```
Requirement already satisfied: python-dateutil>=2.8.2 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from pandas>=0.24->catboost) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from pandas>=0.24->catboost) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from pandas>=0.24->catboost) (2025.2)
Requirement already satisfied: greenlet>=1 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from sqlalchemy>=1.4.2->optuna) (3.0.3)
Requirement already satisfied: MarkupSafe>=0.9.2 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from Mako->alembic>=1.5.0->optuna) (3.0.3)
Requirement already satisfied: contourpy>=1.0.1 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from matplotlib->catboost) (1.3.3)
Requirement already satisfied: cycler>=0.10 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from matplotlib->catboost) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from matplotlib->catboost) (4.60.1)
Requirement already satisfied: kiwisolver>=1.3.1 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from matplotlib->catboost) (1.4.9)
Requirement already satisfied: pillow>=8 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from matplotlib->catboost) (12.0.0)
Requirement already satisfied: pyparsing>=3 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from matplotlib->catboost) (3.2.5)
Requirement already satisfied: narwhals>=1.15.1 in /home/optert/ML-labs-Tert/.venv/lib/python3.11/site-packages (from plotly->catboost) (2.10.2)
Note: you may need to restart the kernel to use updated packages.
```

```
[68]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import optuna
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score, roc_curve
from sklearn.preprocessing import LabelEncoder, OrdinalEncoder, OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
```

```

from sklearn.model_selection import StratifiedKFold
from catboost import CatBoostClassifier
from optuna.samplers import TPESampler
from sklearn.model_selection import cross_val_score, StratifiedKFold
from optuna.samplers import TPESampler
from optuna.pruners import MedianPruner

[ ]: import os
from pathlib import Path

IN_COLAB = "COLAB_GPU" in os.environ or "google.colab" in str(get_ipython())

if not IN_COLAB:
    ROOT = Path.cwd()

    if ROOT.name == "lab02":
        LAB_ROOT = ROOT
    elif (ROOT / "lab02").exists():
        LAB_ROOT = ROOT / "lab02"
    elif ROOT.name == "notebooks":
        LAB_ROOT = ROOT.parent
    else:
        raise FileNotFoundError("lab02")

DATA_RAW = LAB_ROOT / "data" / "raw"
DATA_RAW.mkdir(parents=True, exist_ok=True)

csv_path = DATA_RAW / "telco.csv"

else:
    from urllib.request import urlretrieve

    url = "https://raw.githubusercontent.com/OlegTertychnyi/ML-labs-Tert/main/lab03/data/raw/telco.csv"
    csv_path = Path("/content/telco.csv")

    if not csv_path.exists():
        urlretrieve(url, csv_path)
        print("Dataset downloaded to Colab")

csv_path

[ ]: PosixPath('/home/optert/ML-labs-Tert/lab03/data/raw/telco.csv')

[69]: df = pd.read_csv(csv_path)
df

```

[69]:	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	\
0	7590-VHVEG	Female	0	Yes	No	1	
1	5575-GNVDE	Male	0	No	No	34	
2	3668-QPYBK	Male	0	No	No	2	
3	7795-CFOCW	Male	0	No	No	45	
4	9237-HQITU	Female	0	No	No	2	
...	
7038	6840-RESVB	Male	0	Yes	Yes	24	
7039	2234-XADUH	Female	0	Yes	Yes	72	
7040	4801-JZAZL	Female	0	Yes	Yes	11	
7041	8361-LTMKD	Male	1	Yes	No	4	
7042	3186-AJIEK	Male	0	No	No	66	
\							
PhoneService	MultipleLines	InternetService	OnlineSecurity	\
0	No	No phone service	DSL	No	
1	Yes	No	DSL	Yes	
2	Yes	No	DSL	Yes	
3	No	No phone service	DSL	Yes	
4	Yes	No	Fiber optic	No	
...	
7038	Yes	Yes	DSL	Yes	
7039	Yes	Yes	Fiber optic	No	
7040	No	No phone service	DSL	Yes	
7041	Yes	Yes	Fiber optic	No	
7042	Yes	No	Fiber optic	Yes	
\							
DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	\
0	No	No	No	No	Month-to-month	Month-to-month	
1	Yes	No	No	No	One year	One year	
2	No	No	No	No	Month-to-month	Month-to-month	
3	Yes	Yes	No	No	One year	One year	
4	No	No	No	No	Month-to-month	Month-to-month	
...	
7038	Yes	Yes	Yes	Yes	One year	One year	
7039	Yes	No	Yes	Yes	One year	One year	
7040	No	No	No	No	Month-to-month	Month-to-month	
7041	No	No	No	No	Month-to-month	Month-to-month	
7042	Yes	Yes	Yes	Yes	Two year	Two year	
\							
PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	\
0	Yes	Electronic check	29.85	29.85	
1	No	Mailed check	56.95	1889.5	
2	Yes	Mailed check	53.85	108.15	
3	No	Bank transfer (automatic)	42.30	1840.75	
4	Yes	Electronic check	70.70	151.65	
...	
7038	Yes	Mailed check	84.80	1990.5	

```

7039           Yes    Credit card (automatic)      103.20      7362.9
7040           Yes    Electronic check          29.60      346.45
7041           Yes    Mailed check              74.40      306.6
7042           Yes  Bank transfer (automatic)  105.65     6844.5

      Churn
0        No
1        No
2       Yes
3        No
4       Yes
...
7038      No
7039      No
7040      No
7041     Yes
7042      No

```

[7043 rows x 21 columns]

```
[70]: for col in df.select_dtypes(include=['object']).columns:
    df[col] = df[col].astype(str).str.strip()
df["TotalCharges"] = pd.to_numeric(df["TotalCharges"], errors="coerce")
df["Churn"] = (df["Churn"] == "Yes").astype(int)
df["TotalCharges"] = df["TotalCharges"].fillna(df["TotalCharges"].median())
```

```
[73]: TARGET = "Churn"
ID_COLS = ["customerID"] if "customerID" in df.columns else []
X = df.drop(columns=[TARGET] + ID_COLS)
y = df[TARGET]

num_cols = X.select_dtypes(include=[np.number]).columns.tolist()
cat_cols = X.select_dtypes(include=["object"]).columns.tolist()

num_cols, cat_cols
```

```
[73]: (['SeniorCitizen', 'tenure', 'MonthlyCharges', 'TotalCharges'],
      ['gender',
       'Partner',
       'Dependents',
       'PhoneService',
       'MultipleLines',
       'InternetService',
       'OnlineSecurity',
       'OnlineBackup',
       'DeviceProtection',
```

```
'TechSupport',
'TreamingTV',
'TreamingMovies',
'Contract',
'PaperlessBilling',
'PaymentMethod'])
```

```
[74]: #
num_pipe = Pipeline([
    ("impute", SimpleImputer(strategy="median")),
    ("scale", StandardScaler())
])

#
cat_pipe = Pipeline([
    ("impute", SimpleImputer(strategy="most_frequent")),
    ("encode", OneHotEncoder(handle_unknown="ignore"))
])

#
preprocessor = ColumnTransformer([
    ("num", num_pipe, num_cols),
    ("cat", cat_pipe, cat_cols)
])

#
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

#
models = {
    "Logistic Regression": LogisticRegression(max_iter=2000, random_state=42),
    "SVM (RBF)": SVC(kernel="rbf", probability=True, random_state=42),
    "KNN (k=7)": KNeighborsClassifier(n_neighbors=7),
    "Random Forest": RandomForestClassifier(n_estimators=500, random_state=42)
}
```

```
[75]: #
pipelines = {
    model_name: Pipeline([
        ("prep", preprocessor),
        ("clf", model)
    ])
    for model_name, model in models.items()
}
```

```

#
results = []

for model_name, pipe in pipelines.items():
    pipe.fit(X_train, y_train)

    prob = pipe.predict_proba(X_test)[:, 1]
    pred = pipe.predict(X_test)

    results.append({
        "Model": model_name,
        "Accuracy": accuracy_score(y_test, pred),
        "Precision": precision_score(y_test, pred),
        "Recall": recall_score(y_test, pred),
        "F1": f1_score(y_test, pred),
        "ROC-AUC": roc_auc_score(y_test, prob)
    })

results_df = pd.DataFrame(results)
results_df

```

	Model	Accuracy	Precision	Recall	F1	ROC-AUC
0	Logistic Regression	0.805536	0.657233	0.558824	0.604046	0.841874
1	SVM (RBF)	0.791341	0.641844	0.483957	0.551829	0.790491
2	KNN (k=7)	0.764372	0.554404	0.572193	0.563158	0.804810
3	Random Forest	0.776437	0.600683	0.470588	0.527736	0.819121

```

[77]: # CatBoost:          OneHot,

cat_features_idx = [X.columns.get_loc(c) for c in cat_cols]

cb_model = CatBoostClassifier(
    iterations=500,
    loss_function="Logloss",
    eval_metric="AUC",
    random_seed=42,
    verbose=False,
    cat_features=cat_features_idx
)

cb_model.fit(X_train, y_train)

y_pred = cb_model.predict(X_test)
y_proba = cb_model.predict_proba(X_test)[:, 1]

results.append({
    "Model": "CatBoost",

```

```

        "Accuracy": accuracy_score(y_test, y_pred),
        "Precision": precision_score(y_test, y_pred),
        "Recall": recall_score(y_test, y_pred),
        "F1": f1_score(y_test, y_pred),
        "ROC-AUC": roc_auc_score(y_test, y_proba)
    })
pd.DataFrame(results)

```

```
[77]:          Model  Accuracy  Precision   Recall      F1  ROC-AUC
0  Logistic Regression  0.805536  0.657233  0.558824  0.604046  0.841874
1            SVM (RBF)  0.791341  0.641844  0.483957  0.551829  0.790491
2           KNN (k=7)  0.764372  0.554404  0.572193  0.563158  0.804810
3       Random Forest  0.776437  0.600683  0.470588  0.527736  0.819121
4         CatBoost  0.805536  0.676056  0.513369  0.583587  0.844270

```

```
[78]: def build_logreg(C, penalty, l1_ratio=None):
    #           solver
    if penalty == "l1":
        solver = "liblinear"
        params = {"penalty": "l1", "C": C}
    elif penalty == "l2":
        solver = "lbfgs"
        params = {"penalty": "l2", "C": C}
    else: # elasticnet
        solver = "saga"
        params = {"penalty": "elasticnet", "C": C, "l1_ratio": l1_ratio}

    model = LogisticRegression(
        solver=solver,
        max_iter=5000,
        random_state=42,
        **params
    )

    return Pipeline([
        ("prep", preprocessor),
        ("clf", model)
    ])

def objective(trial: optuna.Trial) -> float:
    penalty = trial.suggest_categorical("penalty", ["l1", "l2", "elasticnet"])
    C = trial.suggest_float("C", 1e-3, 1e2, log=True)

    l1_ratio = None
    if penalty == "elasticnet":

```

```

    l1_ratio = trial.suggest_float("l1_ratio", 0.0, 1.0)

    pipe = build_logreg(C=C, penalty=penalty, l1_ratio=l1_ratio)

    cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
    score = cross_val_score(pipe, X_train, y_train, cv=cv,
                            scoring="roc_auc", n_jobs=-1)

    return score.mean()

study_lr = optuna.create_study(direction="maximize",
                                sampler=TPESampler(seed=42))
study_lr.optimize(objective, n_trials=30, show_progress_bar=True)

print("Best AUC:", study_lr.best_value)
print("Best Params:", study_lr.best_params)

```

[I 2025-11-21 14:26:53,189] A new study created in memory with name: no-name-13547912-ae20-4d86-89af-1b432b48a319

0%		0/30 [00:00<?, ?it/s]
----	--	-----------------------

[I 2025-11-21 14:26:59,848] Trial 0 finished with value: 0.8461231646197689 and parameters: {'penalty': 'l2', 'C': 0.9846738873614566}. Best is trial 0 with value: 0.8461231646197689.

[I 2025-11-21 14:27:03,789] Trial 1 finished with value: 0.846254108902029 and parameters: {'penalty': 'l1', 'C': 21.42302175774105}. Best is trial 1 with value: 0.846254108902029.

[I 2025-11-21 14:27:04,309] Trial 2 finished with value: 0.8462928493573113 and parameters: {'penalty': 'l2', 'C': 70.72114131472235}. Best is trial 2 with value: 0.8462928493573113.

[I 2025-11-21 14:27:04,756] Trial 3 finished with value: 0.8213951185326644 and parameters: {'penalty': 'l1', 'C': 0.008260808399079604}. Best is trial 2 with value: 0.8462928493573113.

[I 2025-11-21 14:27:05,209] Trial 4 finished with value: 0.8445429449751904 and parameters: {'penalty': 'l2', 'C': 0.028585493941961918}. Best is trial 2 with value: 0.8462928493573113.

[I 2025-11-21 14:27:05,751] Trial 5 finished with value: 0.8443335551539185 and parameters: {'penalty': 'l1', 'C': 0.06789053271698488}. Best is trial 2 with value: 0.8462928493573113.

[I 2025-11-21 14:27:06,281] Trial 6 finished with value: 0.8459348212817492 and parameters: {'penalty': 'l2', 'C': 0.37253938395788866}. Best is trial 2 with value: 0.8462928493573113.

[I 2025-11-21 14:27:07,094] Trial 7 finished with value: 0.8428136137978578 and parameters: {'penalty': 'elasticnet', 'C': 0.007122305833333868, 'l1_ratio': 0.06505159298527952}. Best is trial 2 with value: 0.8462928493573113.

[I 2025-11-21 14:27:07,559] Trial 8 finished with value: 0.8446552990356977 and

```
parameters: {'penalty': 'l2', 'C': 0.03334792728637585}. Best is trial 2 with
value: 0.8462928493573113.
[I 2025-11-21 14:27:07,970] Trial 9 finished with value: 0.8420773194976992 and
parameters: {'penalty': 'l2', 'C': 0.00407559644007287}. Best is trial 2 with
value: 0.8462928493573113.
[I 2025-11-21 14:27:22,317] Trial 10 finished with value: 0.8462855826524814 and
parameters: {'penalty': 'elasticnet', 'C': 46.08678541950922, 'l1_ratio':
0.9695371107878658}. Best is trial 2 with value: 0.8462928493573113.
[I 2025-11-21 14:27:43,287] Trial 11 finished with value: 0.8462831727908314 and
parameters: {'penalty': 'elasticnet', 'C': 89.83132016522954, 'l1_ratio':
0.9076647952825176}. Best is trial 2 with value: 0.8462928493573113.
[I 2025-11-21 14:28:15,043] Trial 12 finished with value: 0.8462993678035176 and
parameters: {'penalty': 'elasticnet', 'C': 11.260271629505922, 'l1_ratio':
0.9702152277517139}. Best is trial 12 with value: 0.8462993678035176.
[I 2025-11-21 14:28:51,321] Trial 13 finished with value: 0.84628400798812 and
parameters: {'penalty': 'elasticnet', 'C': 6.2892520107372265, 'l1_ratio':
0.5585247307792902}. Best is trial 12 with value: 0.8462993678035176.
[I 2025-11-21 14:29:23,097] Trial 14 finished with value: 0.8462759305069749 and
parameters: {'penalty': 'elasticnet', 'C': 5.390363922402939, 'l1_ratio':
0.5780557316801694}. Best is trial 12 with value: 0.8462993678035176.
[I 2025-11-21 14:29:23,615] Trial 15 finished with value: 0.8462969266830216 and
parameters: {'penalty': 'l2', 'C': 6.831117665725749}. Best is trial 12 with
value: 0.8462993678035176.
[I 2025-11-21 14:29:39,632] Trial 16 finished with value: 0.8462355304023438 and
parameters: {'penalty': 'elasticnet', 'C': 2.3880566939107437, 'l1_ratio':
0.18318389605254987}. Best is trial 12 with value: 0.8462993678035176.
[I 2025-11-21 14:30:27,759] Trial 17 finished with value: 0.8463058247088714 and
parameters: {'penalty': 'elasticnet', 'C': 10.172075441673128, 'l1_ratio':
0.7593033273862684}. Best is trial 17 with value: 0.8463058247088714.
[I 2025-11-21 14:30:37,784] Trial 18 finished with value: 0.8460713882489129 and
parameters: {'penalty': 'elasticnet', 'C': 0.49899998666091444, 'l1_ratio':
0.7752033459935036}. Best is trial 17 with value: 0.8463058247088714.
[I 2025-11-21 14:31:54,699] Trial 19 finished with value: 0.8462718795559157 and
parameters: {'penalty': 'elasticnet', 'C': 17.962491814149992, 'l1_ratio':
0.7395095724659485}. Best is trial 17 with value: 0.8463058247088714.
[I 2025-11-21 14:32:11,983] Trial 20 finished with value: 0.8460949144377986 and
parameters: {'penalty': 'elasticnet', 'C': 1.7392244271251245, 'l1_ratio':
0.7613756157041719}. Best is trial 17 with value: 0.8463058247088714.
[I 2025-11-21 14:32:34,249] Trial 21 finished with value: 0.8462775501059274 and
parameters: {'penalty': 'elasticnet', 'C': 8.469553977729813, 'l1_ratio':
0.9936291448929787}. Best is trial 17 with value: 0.8463058247088714.
[I 2025-11-21 14:32:34,660] Trial 22 finished with value: 0.8393614728199157 and
parameters: {'penalty': 'l2', 'C': 0.0011633263858100054}. Best is trial 17 with
value: 0.8463058247088714.
[I 2025-11-21 14:32:38,661] Trial 23 finished with value: 0.8462880540549845 and
parameters: {'penalty': 'l1', 'C': 19.540027649806603}. Best is trial 17 with
value: 0.8463058247088714.
[I 2025-11-21 14:32:59,269] Trial 24 finished with value: 0.846292113797594 and
```

```

parameters: {'penalty': 'elasticnet', 'C': 3.125059166612226, 'l1_ratio': 0.35254572768281606}. Best is trial 17 with value: 0.8463058247088714.
[I 2025-11-21 14:32:59,771] Trial 25 finished with value: 0.8454498441951669 and
parameters: {'penalty': 'l2', 'C': 0.09997778181293299}. Best is trial 17 with
value: 0.8463058247088714.
[I 2025-11-21 14:33:49,327] Trial 26 finished with value: 0.8462985667955918 and
parameters: {'penalty': 'elasticnet', 'C': 13.457197497444264, 'l1_ratio': 0.8448345378628919}. Best is trial 17 with value: 0.8463058247088714.
[I 2025-11-21 14:34:03,045] Trial 27 finished with value: 0.846070631198739 and
parameters: {'penalty': 'elasticnet', 'C': 0.891474572640944, 'l1_ratio': 0.8296001148801442}. Best is trial 17 with value: 0.8463058247088714.
[I 2025-11-21 14:34:36,533] Trial 28 finished with value: 0.846287218857696 and
parameters: {'penalty': 'elasticnet', 'C': 28.753580717122208, 'l1_ratio': 0.6496777263540325}. Best is trial 17 with value: 0.8463058247088714.
[I 2025-11-21 14:34:51,862] Trial 29 finished with value: 0.8460690184376591 and
parameters: {'penalty': 'elasticnet', 'C': 0.9199988595892168, 'l1_ratio': 0.8493385231889122}. Best is trial 17 with value: 0.8463058247088714.
Best AUC: 0.8463058247088714
Best Params: {'penalty': 'elasticnet', 'C': 10.172075441673128, 'l1_ratio': 0.7593033273862684}

```

```
[79]: # Optuna
best_lr_pipe = build_logreg(
    C=study_lr.best_params["C"],
    penalty=study_lr.best_params["penalty"],
    l1_ratio=study_lr.best_params.get("l1_ratio")
)

best_lr_pipe.fit(X_train, y_train)

#
lr_pred = best_lr_pipe.predict(X_test)
lr_proba = best_lr_pipe.predict_proba(X_test)[:, 1]

#
results.append({
    "Model": "LogReg (tuned)",
    "Accuracy": accuracy_score(y_test, lr_pred),
    "Precision": precision_score(y_test, lr_pred),
    "Recall": recall_score(y_test, lr_pred),
    "F1": f1_score(y_test, lr_pred),
    "ROC-AUC": roc_auc_score(y_test, lr_proba)
})
```

```
[80]: pd.DataFrame(results)
```

```
[80]:
```

	Model	Accuracy	Precision	Recall	F1	ROC-AUC
0	Logistic Regression	0.805536	0.657233	0.558824	0.604046	0.841874
1	SVM (RBF)	0.791341	0.641844	0.483957	0.551829	0.790491
2	KNN (k=7)	0.764372	0.554404	0.572193	0.563158	0.804810
3	Random Forest	0.776437	0.600683	0.470588	0.527736	0.819121
4	CatBoost	0.805536	0.676056	0.513369	0.583587	0.844270
5	LogReg (tuned)	0.804826	0.655172	0.558824	0.603175	0.841169

```
[85]: #          catboost
cb_preprocessor = ColumnTransformer(
    transformers=[
        ('num', 'passthrough', num_cols),
        ('cat', 'passthrough', cat_cols),
    ],
    remainder='drop'
)

#          ColumnTransformer(passthrough)
cat_idx = list(range(len(num_cols), len(num_cols) + len(cat_cols)))

def objective_cb(trial: optuna.Trial) -> float:

    params = {
        "learning_rate": trial.suggest_float("learning_rate", 0.03, 0.2, log=True),
        "depth": trial.suggest_int("depth", 4, 6),
        "l2_leaf_reg": trial.suggest_float("l2_leaf_reg", 2.0, 8.0),
        "n_estimators": trial.suggest_int("n_estimators", 300, 1200),
        "bagging_temperature": trial.suggest_float("bagging_temperature", 0.0, 1.0),
        "random_strength": trial.suggest_float("random_strength", 0.2, 1.5),
    }

    cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
    oof_scores = []

    for train_idx, valid_idx in cv.split(X_train, y_train):

        X_tr, X_va = X_train.iloc[train_idx], X_train.iloc[valid_idx]
        y_tr, y_va = y_train.iloc[train_idx], y_train.iloc[valid_idx]

        #
        pre = cb_preprocessor
        pre.fit(X_tr)
        X_tr_t = pre.transform(X_tr)
```

```

X_va_t = pre.transform(X_va)

model = CatBoostClassifier(
    loss_function="Logloss",
    eval_metric="AUC",
    random_seed=42,
    verbose=False,
    cat_features=cat_idx,
    **params
)

model.fit(
    X_tr_t, y_tr,
    eval_set=(X_va_t, y_va),
    early_stopping_rounds=100
)

y_va_proba = model.predict_proba(X_va_t)[:, 1]
oof_scores.append(roc_auc_score(y_va, y_va_proba))

trial.report(float(np.mean(oof_scores)), step=len(oof_scores))
if trial.should_prune():
    raise optuna.TrialPruned()

return float(np.mean(oof_scores))

study_cb = optuna.create_study(
    direction="maximize",
    sampler=TPESampler(seed=42),
    pruner=MedianPruner(n_startup_trials=10)
)

study_cb.optimize(objective_cb, n_trials=25, show_progress_bar=True)

print("Best CB AUC (CV):", study_cb.best_value)
print("Best CB params:", study_cb.best_params)

```

[I 2025-11-21 15:04:34,746] A new study created in memory with name: no-name-72f32a14-ad10-4a04-9e54-2b5195afae05

0% | 0/25 [00:00<?, ?it/s]

[I 2025-11-21 15:04:58,578] Trial 0 finished with value: 0.8490689059448876 and parameters: {'learning_rate': 0.061053157936408836, 'depth': 6, 'l2_leaf_reg': 6.391963650868431, 'n_estimators': 839, 'bagging_temperature': 0.15601864044243652, 'random_strength': 0.40279287643706346}. Best is trial 0 with value: 0.8490689059448876.

[I 2025-11-21 15:05:32,586] Trial 1 finished with value: 0.8500720770375819 and parameters: {'learning_rate': 0.033494758473134406, 'depth': 6, 'l2_leaf_reg': 5.606690070459253, 'n_estimators': 937, 'bagging_temperature': 0.020584494295802447, 'random_strength': 1.4608828078105927}. Best is trial 1 with value: 0.8500720770375819.

[I 2025-11-21 15:05:44,840] Trial 2 finished with value: 0.8500004473922319 and parameters: {'learning_rate': 0.14553854298843344, 'depth': 4, 'l2_leaf_reg': 3.0909498032426037, 'n_estimators': 465, 'bagging_temperature': 0.3042422429595377, 'random_strength': 0.8821833611219092}. Best is trial 1 with value: 0.8500720770375819.

[I 2025-11-21 15:06:04,039] Trial 3 finished with value: 0.8504845760941999 and parameters: {'learning_rate': 0.06807764614983679, 'depth': 4, 'l2_leaf_reg': 5.671117368334277, 'n_estimators': 425, 'bagging_temperature': 0.29214464853521815, 'random_strength': 0.6762703962817993}. Best is trial 3 with value: 0.8504845760941999.

[I 2025-11-21 15:06:24,939] Trial 4 finished with value: 0.8490973075368927 and parameters: {'learning_rate': 0.07126582239713602, 'depth': 6, 'l2_leaf_reg': 3.1980426929501586, 'n_estimators': 763, 'bagging_temperature': 0.5924145688620425, 'random_strength': 0.26038553653599705}. Best is trial 3 with value: 0.8504845760941999.

[I 2025-11-21 15:06:39,523] Trial 5 finished with value: 0.8512647006689061 and parameters: {'learning_rate': 0.09499106760020673, 'depth': 4, 'l2_leaf_reg': 2.390309557911677, 'n_estimators': 1154, 'bagging_temperature': 0.9656320330745594, 'random_strength': 1.2509165525513994}. Best is trial 5 with value: 0.8512647006689061.

[I 2025-11-21 15:07:01,412] Trial 6 finished with value: 0.8502828105008076 and parameters: {'learning_rate': 0.053468155416899, 'depth': 4, 'l2_leaf_reg': 6.105398159072942, 'n_estimators': 696, 'bagging_temperature': 0.12203823484477883, 'random_strength': 0.8437299831446512}. Best is trial 5 with value: 0.8512647006689061.

[I 2025-11-21 15:07:34,153] Trial 7 finished with value: 0.8498331559100777 and parameters: {'learning_rate': 0.032022428000607145, 'depth': 6, 'l2_leaf_reg': 3.5526798896001015, 'n_estimators': 896, 'bagging_temperature': 0.31171107608941095, 'random_strength': 0.876088427531154}. Best is trial 5 with value: 0.8512647006689061.

[I 2025-11-21 15:07:51,579] Trial 8 finished with value: 0.8510134659982196 and parameters: {'learning_rate': 0.08463707239000723, 'depth': 4, 'l2_leaf_reg': 7.817507766587351, 'n_estimators': 998, 'bagging_temperature': 0.9394989415641891, 'random_strength': 1.3632755555559435}. Best is trial 5 with value: 0.8512647006689061.

[I 2025-11-21 15:08:08,957] Trial 9 finished with value: 0.8484031447883756 and parameters: {'learning_rate': 0.09326877515737107, 'depth': 6, 'l2_leaf_reg': 2.5309550123115168, 'n_estimators': 476, 'bagging_temperature': 0.045227288910538066, 'random_strength': 0.6229294299922437}. Best is trial 5 with value: 0.8512647006689061.

[I 2025-11-21 15:08:18,694] Trial 10 pruned.

[I 2025-11-21 15:08:29,158] Trial 11 pruned.

[I 2025-11-21 15:08:46,464] Trial 12 finished with value: 0.851103143719714 and

```

parameters: {'learning_rate': 0.11250097158764306, 'depth': 4, 'l2_leaf_reg': 7.747857049445863, 'n_estimators': 1055, 'bagging_temperature': 0.7617805327747955, 'random_strength': 1.1856300139875224}. Best is trial 5 with value: 0.8512647006689061.

[I 2025-11-21 15:09:01,821] Trial 13 finished with value: 0.8505059288164574 and parameters: {'learning_rate': 0.1259013427444294, 'depth': 5, 'l2_leaf_reg': 4.437934368576814, 'n_estimators': 1072, 'bagging_temperature': 0.760596035264475, 'random_strength': 1.1089110643063502}. Best is trial 5 with value: 0.8512647006689061.

[I 2025-11-21 15:09:13,874] Trial 14 finished with value: 0.8509299785050499 and parameters: {'learning_rate': 0.1670917479157165, 'depth': 4, 'l2_leaf_reg': 2.1190426987904676, 'n_estimators': 1086, 'bagging_temperature': 0.7371824291629284, 'random_strength': 1.077988831634637}. Best is trial 5 with value: 0.8512647006689061.

[I 2025-11-21 15:09:28,855] Trial 15 finished with value: 0.8516166308274229 and parameters: {'learning_rate': 0.0968412457373984, 'depth': 4, 'l2_leaf_reg': 6.888702314029387, 'n_estimators': 618, 'bagging_temperature': 0.8095789183531144, 'random_strength': 1.4964672476949448}. Best is trial 15 with value: 0.8516166308274229.

[I 2025-11-21 15:09:52,598] Trial 16 finished with value: 0.8506173260632292 and parameters: {'learning_rate': 0.05244335214874761, 'depth': 5, 'l2_leaf_reg': 6.8483962066829855, 'n_estimators': 625, 'bagging_temperature': 0.8467638526830317, 'random_strength': 1.4625324456243842}. Best is trial 15 with value: 0.8516166308274229.

[I 2025-11-21 15:10:08,168] Trial 17 finished with value: 0.8516926240122913 and parameters: {'learning_rate': 0.0934604919339652, 'depth': 4, 'l2_leaf_reg': 4.84634432283171, 'n_estimators': 646, 'bagging_temperature': 0.5795660494660647, 'random_strength': 1.3335184557179816}. Best is trial 17 with value: 0.8516926240122913.

[I 2025-11-21 15:10:11,012] Trial 18 pruned.
[I 2025-11-21 15:10:15,029] Trial 19 pruned.
[I 2025-11-21 15:10:32,108] Trial 20 finished with value: 0.8512299252028568 and parameters: {'learning_rate': 0.08084925472114562, 'depth': 4, 'l2_leaf_reg': 5.126041399127765, 'n_estimators': 586, 'bagging_temperature': 0.4578028138312887, 'random_strength': 1.4113618661911014}. Best is trial 17 with value: 0.8516926240122913.

[I 2025-11-21 15:10:34,746] Trial 21 pruned.
[I 2025-11-21 15:10:45,078] Trial 22 pruned.
[I 2025-11-21 15:10:47,098] Trial 23 pruned.
[I 2025-11-21 15:10:49,773] Trial 24 pruned.

Best CB AUC (CV): 0.8516926240122913

Best CB params: {'learning_rate': 0.0934604919339652, 'depth': 4, 'l2_leaf_reg': 4.84634432283171, 'n_estimators': 646, 'bagging_temperature': 0.5795660494660647, 'random_strength': 1.3335184557179816}

```

```
[86]: #     results
# results = []
```

```

pre = cb_preprocessor
pre.fit(X_train, y_train)

X_train_t = pre.transform(X_train)
X_test_t = pre.transform(X_test)

best_cb_model = CatBoostClassifier(
    loss_function="Logloss",
    eval_metric="AUC",
    random_seed=42,
    verbose=False,
    cat_features=cat_idx,
    **study_cb.best_params
)

best_cb_model.fit(
    X_train_t, y_train,
    eval_set=(X_test_t, y_test),
    early_stopping_rounds=100
)

y_scores = best_cb_model.predict_proba(X_test_t)[:, 1]
y_pred = best_cb_model.predict(X_test_t)

results_catboost = {
    "Model": "Tuned CatBoost",
    "Accuracy": accuracy_score(y_test, y_pred),
    "Precision": precision_score(y_test, y_pred, zero_division=0),
    "Recall": recall_score(y_test, y_pred, zero_division=0),
    "F1": f1_score(y_test, y_pred, zero_division=0),
    "ROC-AUC": roc_auc_score(y_test, y_scores),
}

results.append(results_catboost)

```

[89]: pd.DataFrame(results)

	Model	Accuracy	Precision	Recall	F1	ROC-AUC
0	Logistic Regression	0.805536	0.657233	0.558824	0.604046	0.841874
1	SVM (RBF)	0.791341	0.641844	0.483957	0.551829	0.790491
2	KNN (k=7)	0.764372	0.554404	0.572193	0.563158	0.804810
3	Random Forest	0.776437	0.600683	0.470588	0.527736	0.819121
4	CatBoost	0.805536	0.676056	0.513369	0.583587	0.844270
5	LogReg (tuned)	0.804826	0.655172	0.558824	0.603175	0.841169
6	Tuned CatBoost	0.800568	0.656566	0.521390	0.581222	0.846425
7	Tuned CatBoost	0.808375	0.673333	0.540107	0.599407	0.846924

```
[90]: def optimize_threshold(trial):
    thr = trial.suggest_float("threshold", 0.05, 0.95)

    y_pred = (y_scores >= thr).astype(int)

    prec = precision_score(y_test, y_pred, zero_division=0)
    rec = recall_score(y_test, y_pred, zero_division=0)
    f1 = f1_score(y_test, y_pred, zero_division=0)

    #      F1-      +
    score = f1 + 0.1 * prec + 0.1 * rec

    return score

study_thr = optuna.create_study(direction="maximize")
study_thr.optimize(optimize_threshold, n_trials=100)

study_thr.best_params
```

[I 2025-11-21 15:12:34,628] A new study created in memory with name: no-name-7a439bee-69bc-403a-95c1-03f88253ac4f
[I 2025-11-21 15:12:34,652] Trial 0 finished with value: 0.7135055558715164 and parameters: {'threshold': 0.11685412729686837}. Best is trial 0 with value: 0.7135055558715164.
[I 2025-11-21 15:12:34,674] Trial 1 finished with value: 0.6601793180421639 and parameters: {'threshold': 0.5577852368603629}. Best is trial 0 with value: 0.7135055558715164.
[I 2025-11-21 15:12:34,697] Trial 2 finished with value: 0.6669626505707628 and parameters: {'threshold': 0.05860960090985782}. Best is trial 0 with value: 0.7135055558715164.
[I 2025-11-21 15:12:34,717] Trial 3 finished with value: 0.1111730572306292 and parameters: {'threshold': 0.9175107050922531}. Best is trial 0 with value: 0.7135055558715164.
[I 2025-11-21 15:12:34,735] Trial 4 finished with value: 0.7220351592757475 and parameters: {'threshold': 0.15395130861953726}. Best is trial 4 with value: 0.7220351592757475.
[I 2025-11-21 15:12:34,753] Trial 5 finished with value: 0.7500976659583882 and parameters: {'threshold': 0.2511200852040709}. Best is trial 5 with value: 0.7500976659583882.
[I 2025-11-21 15:12:34,770] Trial 6 finished with value: 0.7346228323659136 and parameters: {'threshold': 0.18318689434987068}. Best is trial 5 with value: 0.7500976659583882.
[I 2025-11-21 15:12:34,789] Trial 7 finished with value: 0.7597079391197038 and parameters: {'threshold': 0.37339999023427917}. Best is trial 7 with value: 0.7597079391197038.
[I 2025-11-21 15:12:34,807] Trial 8 finished with value: 0.22271786417916448 and

```
parameters: {'threshold': 0.8314907917681222}. Best is trial 7 with value:  
0.7597079391197038.  
[I 2025-11-21 15:12:34,824] Trial 9 finished with value: 0.7542273294648217 and  
parameters: {'threshold': 0.39075371460325997}. Best is trial 7 with value:  
0.7597079391197038.  
[I 2025-11-21 15:12:34,847] Trial 10 finished with value: 0.571718629360075 and  
parameters: {'threshold': 0.6077385047976682}. Best is trial 7 with value:  
0.7597079391197038.  
[I 2025-11-21 15:12:34,870] Trial 11 finished with value: 0.7594850410333979 and  
parameters: {'threshold': 0.3748614287775843}. Best is trial 7 with value:  
0.7597079391197038.  
[I 2025-11-21 15:12:34,895] Trial 12 finished with value: 0.7448246141958109 and  
parameters: {'threshold': 0.40879901978409955}. Best is trial 7 with value:  
0.7597079391197038.  
[I 2025-11-21 15:12:34,917] Trial 13 finished with value: 0.7652491441867983 and  
parameters: {'threshold': 0.33719432709151986}. Best is trial 13 with value:  
0.7652491441867983.  
[I 2025-11-21 15:12:34,939] Trial 14 finished with value: 0.4663168050468765 and  
parameters: {'threshold': 0.6919791403148126}. Best is trial 13 with value:  
0.7652491441867983.  
[I 2025-11-21 15:12:34,961] Trial 15 finished with value: 0.75231290442989 and  
parameters: {'threshold': 0.2993404130092395}. Best is trial 13 with value:  
0.7652491441867983.  
[I 2025-11-21 15:12:34,983] Trial 16 finished with value: 0.721688945308393 and  
parameters: {'threshold': 0.49081510509610454}. Best is trial 13 with value:  
0.7652491441867983.  
[I 2025-11-21 15:12:35,009] Trial 17 finished with value: 0.7516698836165777 and  
parameters: {'threshold': 0.2912478246912531}. Best is trial 13 with value:  
0.7652491441867983.  
[I 2025-11-21 15:12:35,032] Trial 18 finished with value: 0.47098728090238096  
and parameters: {'threshold': 0.686321089270544}. Best is trial 13 with value:  
0.7652491441867983.  
[I 2025-11-21 15:12:35,053] Trial 19 finished with value: 0.7441630850136686 and  
parameters: {'threshold': 0.44038984813741033}. Best is trial 13 with value:  
0.7652491441867983.  
[I 2025-11-21 15:12:35,077] Trial 20 finished with value: 0.7468671855177988 and  
parameters: {'threshold': 0.23566305721371153}. Best is trial 13 with value:  
0.7652491441867983.  
[I 2025-11-21 15:12:35,101] Trial 21 finished with value: 0.7653046884593685 and  
parameters: {'threshold': 0.35212902830297127}. Best is trial 21 with value:  
0.7653046884593685.  
[I 2025-11-21 15:12:35,125] Trial 22 finished with value: 0.7663261044953519 and  
parameters: {'threshold': 0.33277007517503965}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,153] Trial 23 finished with value: 0.758343359197968 and  
parameters: {'threshold': 0.3162293952762993}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,176] Trial 24 finished with value: 0.7091219070062539 and
```

```
parameters: {'threshold': 0.5164034956686969}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,197] Trial 25 finished with value: 0.7218859817317493 and  
parameters: {'threshold': 0.47498893491554534}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,226] Trial 26 finished with value: 0.7419723405017523 and  
parameters: {'threshold': 0.20774425642501462}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,284] Trial 27 finished with value: 0.7575322126844863 and  
parameters: {'threshold': 0.31308599991648783}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,319] Trial 28 finished with value: 0.6276862540128155 and  
parameters: {'threshold': 0.5791299285833059}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,345] Trial 29 finished with value: 0.7081106713880219 and  
parameters: {'threshold': 0.10893617658719207}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,386] Trial 30 finished with value: 0.7588349174915253 and  
parameters: {'threshold': 0.3646126861992251}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,432] Trial 31 finished with value: 0.7584320464988986 and  
parameters: {'threshold': 0.3569024341663302}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,456] Trial 32 finished with value: 0.7421688963708485 and  
parameters: {'threshold': 0.4381749736085038}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,479] Trial 33 finished with value: 0.7515309197775456 and  
parameters: {'threshold': 0.2559951631569415}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,500] Trial 34 finished with value: 0.7640512128428225 and  
parameters: {'threshold': 0.34820357204648295}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,522] Trial 35 finished with value: 0.7081106713880219 and  
parameters: {'threshold': 0.10890485279587214}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,545] Trial 36 finished with value: 0.6798769701902224 and  
parameters: {'threshold': 0.5347858360107894}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,567] Trial 37 finished with value: 0.7227718360071301 and  
parameters: {'threshold': 0.15573633441551182}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,589] Trial 38 finished with value: 0.7400243827848484 and  
parameters: {'threshold': 0.45320327417887357}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,613] Trial 39 finished with value: 0.7516527140352012 and  
parameters: {'threshold': 0.21724917200664828}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,635] Trial 40 finished with value: 0.7521000243971476 and
```

```
parameters: {'threshold': 0.2689335462743563}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,658] Trial 41 finished with value: 0.7630842960351176 and  
parameters: {'threshold': 0.33932577243385964}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,685] Trial 42 finished with value: 0.7597062308892572 and  
parameters: {'threshold': 0.32700354588859903}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,709] Trial 43 finished with value: 0.7429454390451833 and  
parameters: {'threshold': 0.4054694128405141}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,731] Trial 44 finished with value: 0.7663261044953519 and  
parameters: {'threshold': 0.33327410037715666}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,753] Trial 45 finished with value: 0.6596758802553484 and  
parameters: {'threshold': 0.05078703219784453}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,776] Trial 46 finished with value: 0.0 and parameters:  
{'threshold': 0.9392697411344801}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,800] Trial 47 finished with value: 0.7290509003428108 and  
parameters: {'threshold': 0.16927686120778657}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,823] Trial 48 finished with value: 0.7441395138986859 and  
parameters: {'threshold': 0.4035449715676013}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,848] Trial 49 finished with value: 0.7480466743384327 and  
parameters: {'threshold': 0.2800651972529995}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,870] Trial 50 finished with value: 0.7388081434666804 and  
parameters: {'threshold': 0.19856695098845215}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,891] Trial 51 finished with value: 0.7663261044953519 and  
parameters: {'threshold': 0.3333496509809824}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,916] Trial 52 finished with value: 0.760425917828994 and  
parameters: {'threshold': 0.35642686981145577}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,942] Trial 53 finished with value: 0.7452941322300792 and  
parameters: {'threshold': 0.4292306904952275}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,965] Trial 54 finished with value: 0.7479595120700616 and  
parameters: {'threshold': 0.24609106881816672}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:35,987] Trial 55 finished with value: 0.19664956891847651  
and parameters: {'threshold': 0.8526703732605443}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,013] Trial 56 finished with value: 0.7515309505100843 and
```

```
parameters: {'threshold': 0.2985856705224138}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,036] Trial 57 finished with value: 0.760676691748065 and  
parameters: {'threshold': 0.38775692797740013}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,058] Trial 58 finished with value: 0.726066847787673 and  
parameters: {'threshold': 0.4656794632714144}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,079] Trial 59 finished with value: 0.5221964518585295 and  
parameters: {'threshold': 0.649546409854223}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,101] Trial 60 finished with value: 0.7663261044953519 and  
parameters: {'threshold': 0.3337459360930574}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,125] Trial 61 finished with value: 0.7597062308892572 and  
parameters: {'threshold': 0.3268699038242207}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,147] Trial 62 finished with value: 0.7480466743384327 and  
parameters: {'threshold': 0.28067010730218606}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,170] Trial 63 finished with value: 0.7603973027352393 and  
parameters: {'threshold': 0.3756695085159267}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,208] Trial 64 finished with value: 0.7630842960351176 and  
parameters: {'threshold': 0.33973775825226005}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,250] Trial 65 finished with value: 0.7416665047828808 and  
parameters: {'threshold': 0.4238707101697846}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,274] Trial 66 finished with value: 0.7482968065269546 and  
parameters: {'threshold': 0.24369835131142392}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,320] Trial 67 finished with value: 0.7229864062769709 and  
parameters: {'threshold': 0.5014127577528869}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,366] Trial 68 finished with value: 0.7617719691535769 and  
parameters: {'threshold': 0.3556515907506062}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,393] Trial 69 finished with value: 0.7517128340409878 and  
parameters: {'threshold': 0.3038441033600262}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,419] Trial 70 finished with value: 0.7506390615024359 and  
parameters: {'threshold': 0.22099064216062375}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,445] Trial 71 finished with value: 0.7652491441867983 and  
parameters: {'threshold': 0.33552825859169505}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,473] Trial 72 finished with value: 0.7579030537409612 and
```

```
parameters: {'threshold': 0.38565625782058255}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,495] Trial 73 finished with value: 0.7521000243971476 and  
parameters: {'threshold': 0.26872755081500144}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,523] Trial 74 finished with value: 0.758343359197968 and  
parameters: {'threshold': 0.3142575115345008}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,552] Trial 75 finished with value: 0.7402789343701295 and  
parameters: {'threshold': 0.4165975383090226}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,586] Trial 76 finished with value: 0.7652491441867983 and  
parameters: {'threshold': 0.33602024303550265}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,609] Trial 77 finished with value: 0.7146263249905381 and  
parameters: {'threshold': 0.13284575630459156}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,629] Trial 78 finished with value: 0.753215113956802 and  
parameters: {'threshold': 0.29347370052914024}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,649] Trial 79 finished with value: 0.7597079391197038 and  
parameters: {'threshold': 0.3730324527870485}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,668] Trial 80 finished with value: 0.7211073353061184 and  
parameters: {'threshold': 0.48349211437489276}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,694] Trial 81 finished with value: 0.7630842960351176 and  
parameters: {'threshold': 0.34076465498240327}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,726] Trial 82 finished with value: 0.7639247113548626 and  
parameters: {'threshold': 0.34329923090896364}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,778] Trial 83 finished with value: 0.7509715923485218 and  
parameters: {'threshold': 0.39668150183420803}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,829] Trial 84 finished with value: 0.7400243827848484 and  
parameters: {'threshold': 0.4519228149188659}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,887] Trial 85 finished with value: 0.7580717048636881 and  
parameters: {'threshold': 0.3195446883414715}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,968] Trial 86 finished with value: 0.750956440373335 and  
parameters: {'threshold': 0.26595181010352065}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:36,998] Trial 87 finished with value: 0.7585990882407969 and  
parameters: {'threshold': 0.3671700342073786}. Best is trial 22 with value:  
0.7663261044953519.  
[I 2025-11-21 15:12:37,057] Trial 88 finished with value: 0.7325771995820851 and
```

```

parameters: {'threshold': 0.18553420186239084}. Best is trial 22 with value:
0.7663261044953519.
[I 2025-11-21 15:12:37,113] Trial 89 finished with value: 0.7530967276796952 and
parameters: {'threshold': 0.30294927362321766}. Best is trial 22 with value:
0.7663261044953519.
[I 2025-11-21 15:12:37,159] Trial 90 finished with value: 0.7482968065269546 and
parameters: {'threshold': 0.24147109970416575}. Best is trial 22 with value:
0.7663261044953519.
[I 2025-11-21 15:12:37,214] Trial 91 finished with value: 0.7639247113548626 and
parameters: {'threshold': 0.34363573121895685}. Best is trial 22 with value:
0.7663261044953519.
[I 2025-11-21 15:12:37,259] Trial 92 finished with value: 0.7597062308892572 and
parameters: {'threshold': 0.32412769950439513}. Best is trial 22 with value:
0.7663261044953519.
[I 2025-11-21 15:12:37,290] Trial 93 finished with value: 0.7573108669155825 and
parameters: {'threshold': 0.35978505367082625}. Best is trial 22 with value:
0.7663261044953519.
[I 2025-11-21 15:12:37,341] Trial 94 finished with value: 0.7480466743384327 and
parameters: {'threshold': 0.28077416229393953}. Best is trial 22 with value:
0.7663261044953519.
[I 2025-11-21 15:12:37,398] Trial 95 finished with value: 0.7441395138986859 and
parameters: {'threshold': 0.4040676856954322}. Best is trial 22 with value:
0.7663261044953519.
[I 2025-11-21 15:12:37,447] Trial 96 finished with value: 0.7630842960351176 and
parameters: {'threshold': 0.3417713894801406}. Best is trial 22 with value:
0.7663261044953519.
[I 2025-11-21 15:12:37,509] Trial 97 finished with value: 0.7579030537409612 and
parameters: {'threshold': 0.3835512123282778}. Best is trial 22 with value:
0.7663261044953519.
[I 2025-11-21 15:12:37,561] Trial 98 finished with value: 0.7522836198408835 and
parameters: {'threshold': 0.2857802122405756}. Best is trial 22 with value:
0.7663261044953519.
[I 2025-11-21 15:12:37,606] Trial 99 finished with value: 0.758887957710403 and
parameters: {'threshold': 0.32057825203798207}. Best is trial 22 with value:
0.7663261044953519.

```

[90]: {'threshold': 0.33277007517503965}

```

[106]: #
best_thr = study_thr.best_params["threshold"]

y_pred_thr = (y_scores >= best_thr).astype(int)

results_catboost_thr = {
    "Model": "Tuned CatBoost + Tuned Threshold",
    "Accuracy": accuracy_score(y_test, y_pred_thr),
    "Precision": precision_score(y_test, y_pred_thr, zero_division=0),
}

```

```

        "Recall": recall_score(y_test, y_pred_thr, zero_division=0),
        "F1 Score": f1_score(y_test, y_pred_thr, zero_division=0),
        "ROC AUC": roc_auc_score(y_test, y_scores)
    }

results.append(results_catboost_thr)

pd.DataFrame(results)

```

[106]:

	Model	Accuracy	Precision	Recall	F1	\
0	Logistic Regression	0.805536	0.657233	0.558824	0.604046	
1	SVM (RBF)	0.791341	0.641844	0.483957	0.551829	
2	KNN (k=7)	0.764372	0.554404	0.572193	0.563158	
3	Random Forest	0.776437	0.600683	0.470588	0.527736	
4	CatBoost	0.805536	0.676056	0.513369	0.583587	
5	LogReg (tuned)	0.804826	0.655172	0.558824	0.603175	
7	Tuned CatBoost	0.808375	0.673333	0.540107	0.599407	
9	Tuned CatBoost + Tuned Threshold	0.773598	0.554672	0.745989	0.636260	

	ROC-AUC
0	0.841874
1	0.790491
2	0.804810
3	0.819121
4	0.844270
5	0.841169
7	0.846924
9	0.846924

[109]:

```

# 1)           - CatBoost           ,          X_train
feature_names = X_train.columns.tolist()

# 2)
fi = pd.DataFrame({
    "Feature": feature_names,
    "Baseline CatBoost Importance": cb_model.get_feature_importance(),
    "Tuned CatBoost Importance": best_cb_model.get_feature_importance(),
})
fi

```

[109]:

	Feature	Baseline CatBoost Importance	Tuned CatBoost Importance
0	gender	2.263302	0.471446
1	SeniorCitizen	1.348421	13.568037
2	Partner	0.982945	10.696227
3	Dependents	1.622667	10.103402
4	tenure	14.373637	1.121126

5	PhoneService	0.485607	0.116135
6	MultipleLines	4.932329	0.617744
7	InternetService	8.435215	0.085294
8	OnlineSecurity	5.416715	2.485342
9	OnlineBackup	4.107117	11.642309
10	DeviceProtection	1.470028	6.263567
11	TechSupport	3.892558	2.094057
12	StreamingTV	2.887360	0.296493
13	StreamingMovies	2.242390	2.700191
14	Contract	13.756544	0.960772
15	PaperlessBilling	2.697946	1.306437
16	PaymentMethod	6.112204	29.947958
17	MonthlyCharges	11.275231	1.858162
18	TotalCharges	11.697784	3.665302

```
[113]: fig, axes = plt.subplots(1, 2, figsize=(14, 6), sharey=True)
axes[0].barh(fi["Feature"], fi["Tuned CatBoost Importance"], color="royalblue")
axes[0].set_title("Tuned CatBoost")
axes[0].set_xlabel("Importance")

axes[1].barh(fi["Feature"], fi["Baseline CatBoost Importance"], color="orange")
axes[1].set_title("Baseline CatBoost")
axes[1].set_xlabel("Importance")
plt.show()
```

